

Unit 01.03.01

CS 5220:

COMPUTER COMMUNICATIONS

Berkeley Socket API - I


XIAOBO ZHOU, Ph.D.

Professor, Department of Computer Science

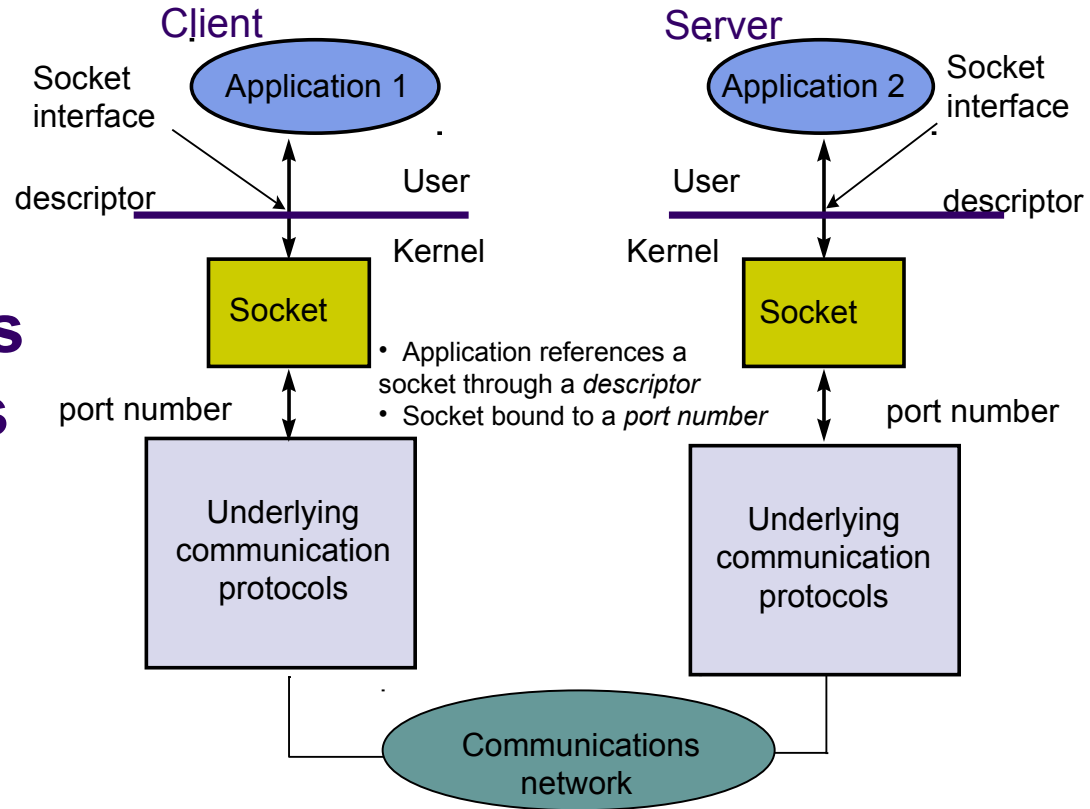


Berkeley Socket API

- Berkeley UNIX Sockets API
 - Abstraction for applications to send & receive data
 - Applications create sockets that “plug into” network
 - Applications write/read to/from sockets
 - Implemented in the kernel
 - Facilitates development of network applications
 - Hides details of underlying protocols & mechanisms
- Also in Windows, Linux, and other OS's



Communications through Sockets



Transport Protocols



- Host computers run two transport protocols on top of IP to enable process-to-process communications
- *User Datagram Protocol* (UDP) enables best-effort connectionless transfer of individual block of information
- *Transmission Control Protocol* (TCP) enables connection-oriented reliable transfer of a stream of bytes
- Two services though Sockets: connection-oriented and connection-less

Stream Mode of Service



Connection-oriented (TCP)


- First, setup connection between two peer application processes
- Then, reliable bidirectional in-sequence transfer of *byte stream* (boundaries not preserved in transfer)
- Multiple write/read between peer processes
- Finally, connection release

Connectionless (UDP)

- Immediate transfer of one block of information (boundaries preserved)
- No setup overhead & delay
- Destination address with each block
- Send/receive to/from multiple peer processes
- Best-effort service only
 - Possible out-of-order
 - Possible loss



Client & Server Differences

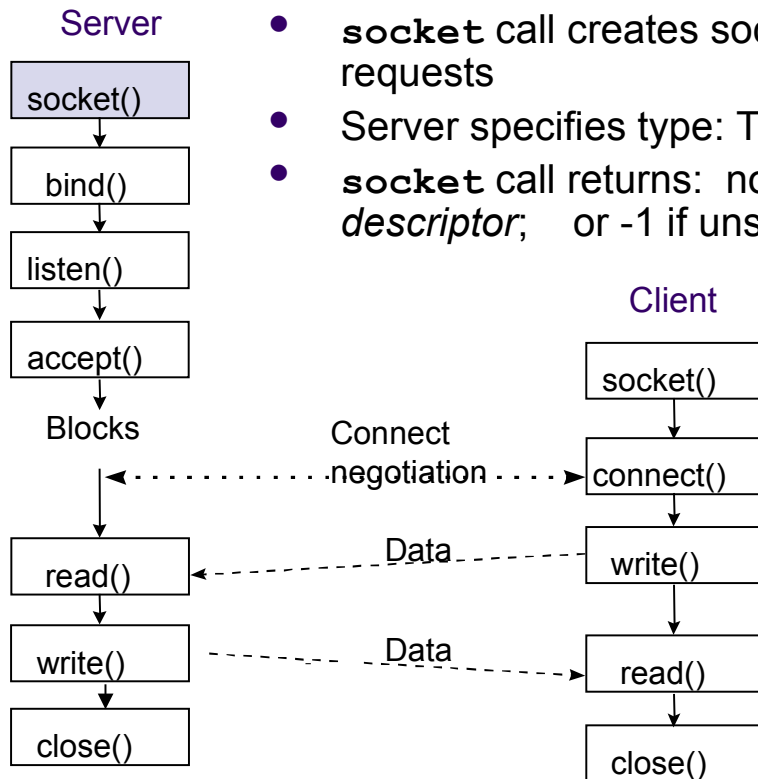


- Server
 - Specifies well-known port # when creating socket
 - May have multiple IP addresses (net interfaces)
 - Waits passively for client requests
- Client
 - Assigned ephemeral port #
 - Initiates communications with server
 - Needs to know server's IP address & port #
 - DNS for URL & server well-known port #
 - Server learns client's address & port #




Server does *Passive Open*

- **socket** call creates socket to *listen* for connection requests
- Server specifies type: TCP (stream)
- **socket** call returns: non-negative integer *descriptor*; or -1 if unsuccessful




Socket Calls for Connection-Oriented Mode



- **bind** assigns local address & port # to socket with specified descriptor
- Can wildcard IP address for multiple net interfaces
- **bind** call returns: 0 (success); or -1 (failure)
- Failure if port # already in use or if reuse option not set

Client




Connect
negotiation


Data

Data

Socket Calls for Connection-Oriented Mode



- **listen** indicates to TCP readiness to receive connection requests for socket with given descriptor
- Parameter specifies max number of requests that may be queued while waiting for server to accept them
- **listen** call returns: 0 (success); or -1 (failure)



Connect
negotiation

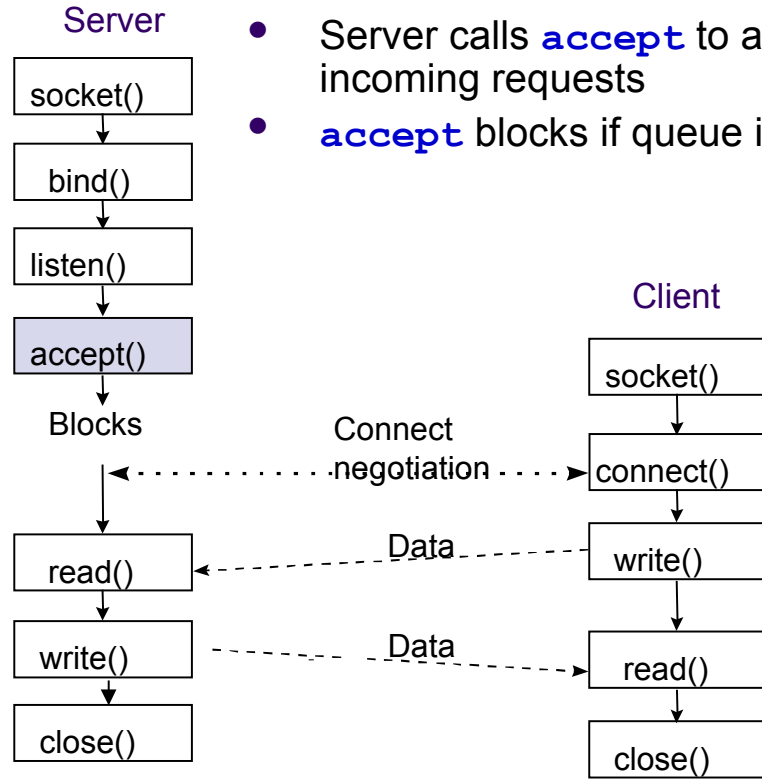

Data

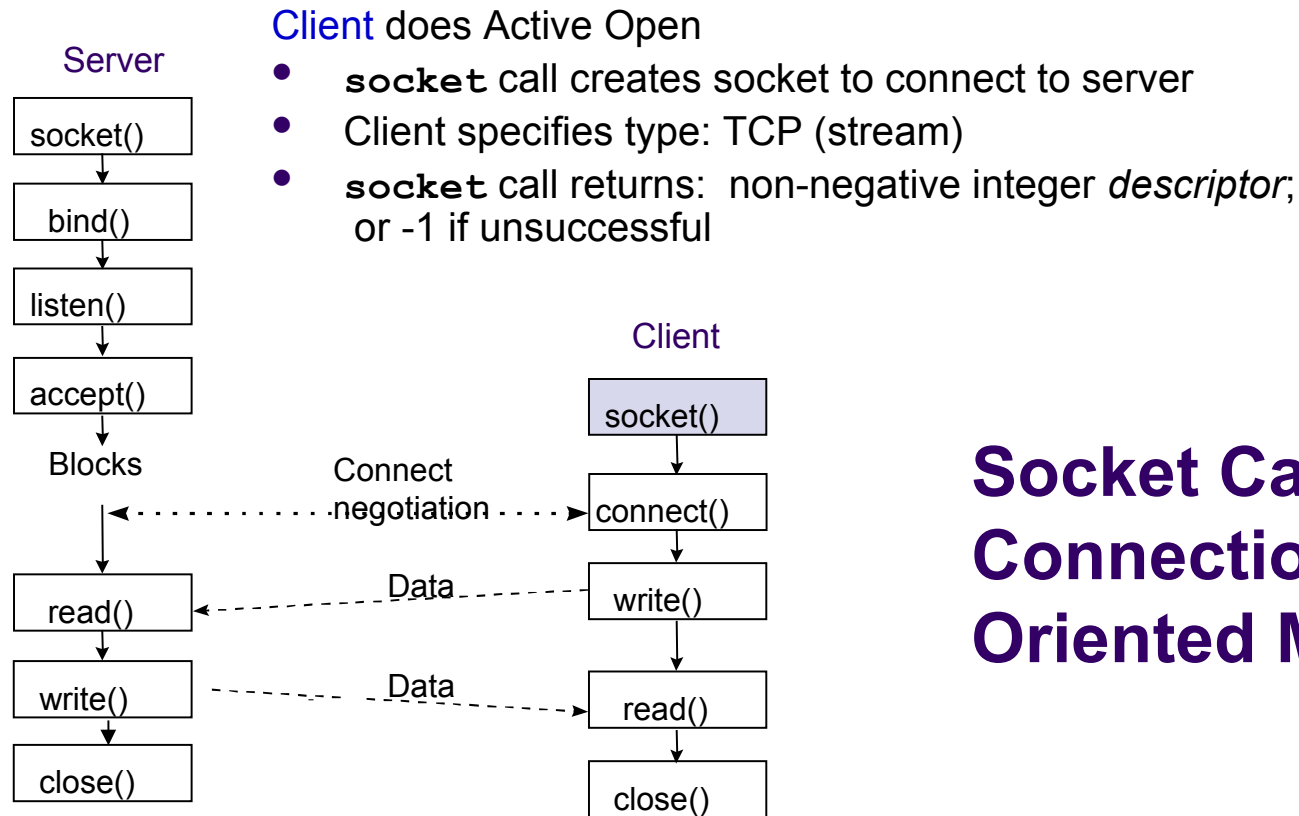
Data

Socket Calls for Connection-Oriented Mode




- Server calls **accept** to accept incoming requests
- **accept** blocks if queue is empty






Socket Calls for Connection-Oriented Mode



- **connect** establishes a connection on the local socket with the specified descriptor to the specified remote address and port #
- **connect** returns 0 if successful; -1 if unsuccessful




Note: **connect** initiates TCP three-way handshake


Connect negotiation

Data

Data



- **accept** wakes with incoming connection request
- **accept** fills client address & port # into address structure
- **accept** call returns: *descriptor of new connection socket* (success); or -1 (failure)
- Client & server use new socket for data transfer
- Original socket continues to listen for new requests




Socket Calls for Connection-Oriented Mode




Data Transfer

- Client or server call **write** to transmit data into a connected socket
- **write** call returns: # bytes transferred (success); or -1 (failure); blocks until all data transferred

Server



Client



Connect
negotiation

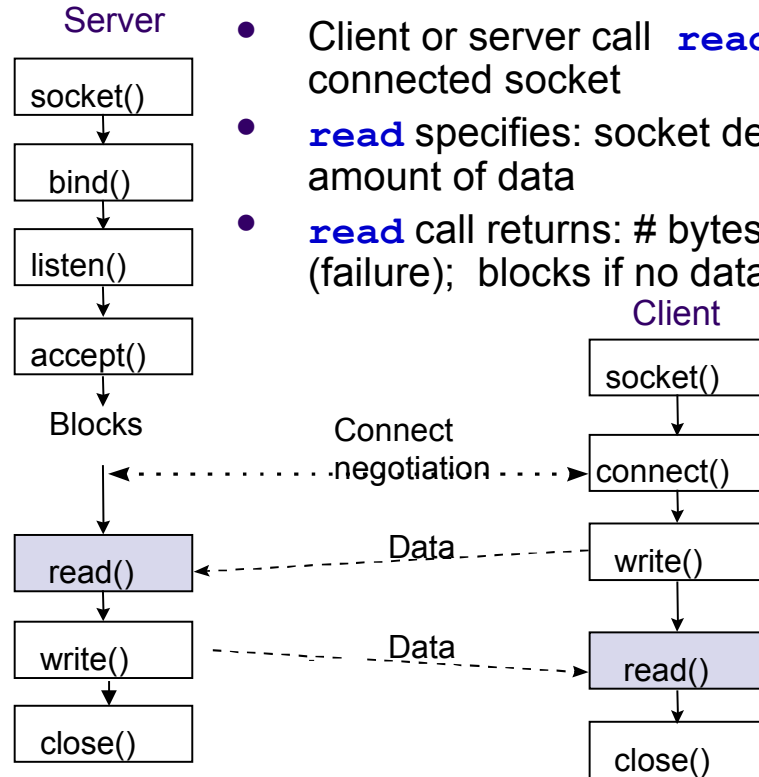
Data

Data



Data Transfer

- Client or server call **read** to receive data from a connected socket
- **read** specifies: socket descriptor; pointer to a buffer; amount of data
- **read** call returns: # bytes read (success); or -1 (failure); blocks if no data arrives

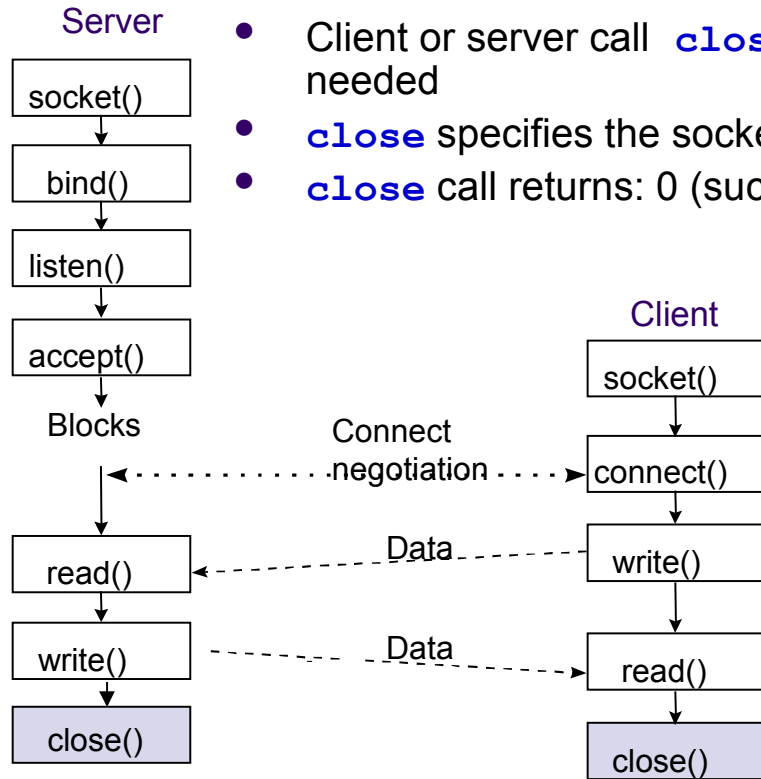


Note: **write** and **read** can be called multiple times to transfer byte streams in both directions



Connection Termination

- Client or server call **close** when socket is no longer needed
- **close** specifies the socket descriptor
- **close** call returns: 0 (success); or -1 (failure)



Note: **close** initiates TCP graceful close sequence

Socket Calls for Connection-Oriented Mode



Summary of the Lesson

- Socket API hides details of underlying protocols & mechanisms