

Unit 01.03.02
CS 5220:
COMPUTER COMMUNICATIONS

Berkeley Socket API - II

XIAOBO ZHOU, Ph.D.


Professor, Department of Computer Science



Stream Mode of Service

Connectionless (UDP)

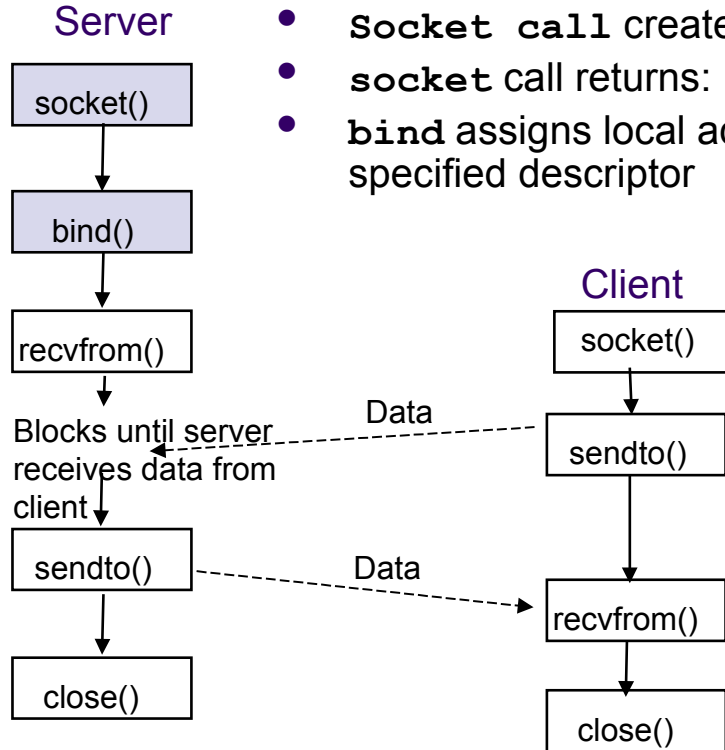
- Immediate transfer of one block of information (boundaries preserved)
- No setup overhead & delay
- Destination address with each block
- Send/receive to/from multiple peer processes
- Best-effort service only
 - Possible out-of-order
 - Possible loss





Server starts first

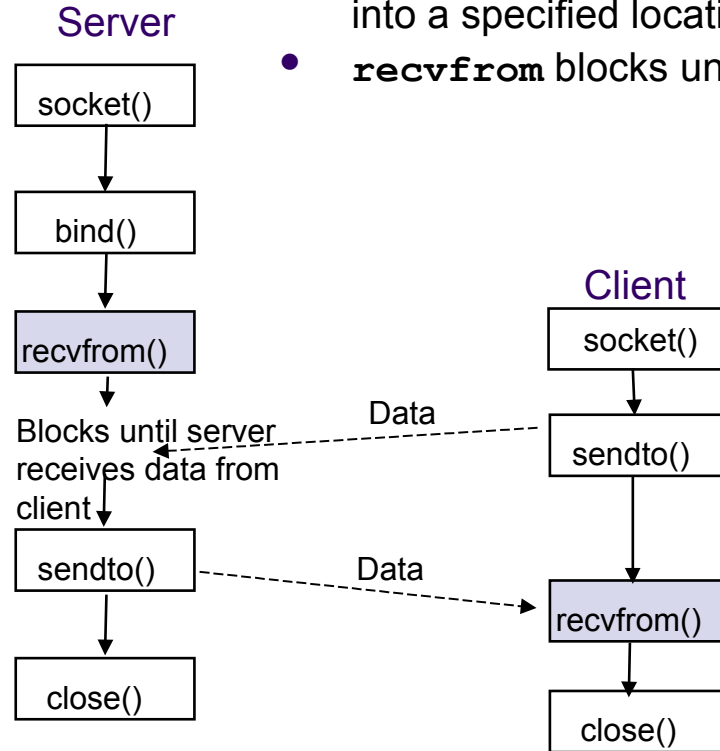
- `socket` call creates socket of type UDP (datagram)
- `socket` call returns: *descriptor*, or -1 if unsuccessful
- `bind` assigns local address & port # to socket with specified descriptor



Socket Calls for Connection-Less Mode



- **recvfrom** copies bytes received in specified socket into a specified location
- **recvfrom** blocks until data arrives

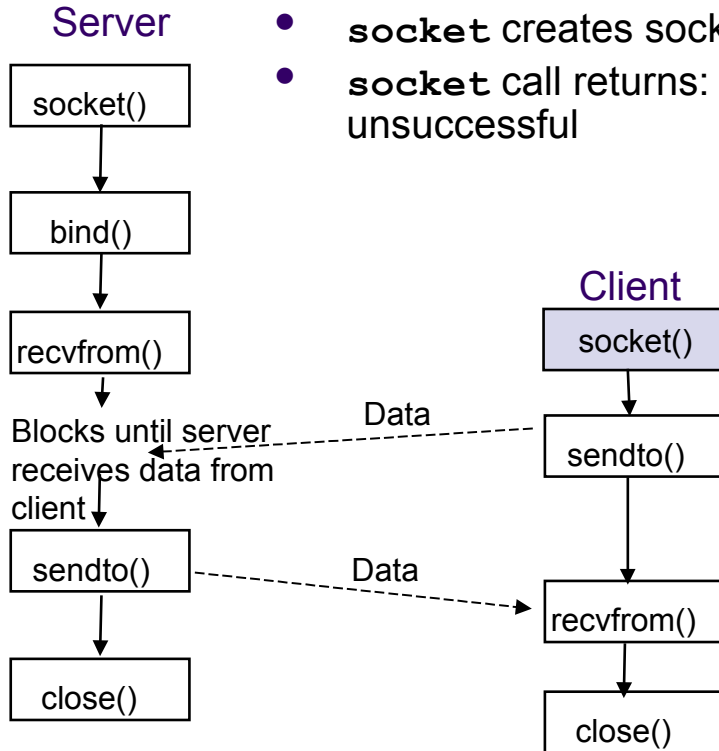


Socket Calls for Connection-Less Mode




Client started


- `socket` creates socket of type UDP (datagram)
- `socket` call returns: *descriptor*; or -1 if unsuccessful



Socket Calls for Connection-Less Mode




Server



- **sendto** transfer bytes in buffer to specified socket
- **sendto** specifies: socket descriptor; pointer to a buffer; amount of data; flags to control transmission behavior; destination address & port #; length of destination address structure
- **sendto** returns: # bytes sent; or -1 if unsuccessful

Client



Blocks until server receives data from client


Data

Data




- **recvfrom** wakes when data arrives
- **recvfrom** specifies: socket descriptor; pointer to a buffer to put data; max # bytes to put in buffer; control flags; copies: sender address & port #; length of sender address structure
- **recvfrom** returns # bytes received or -1 (failure)

Server



Client

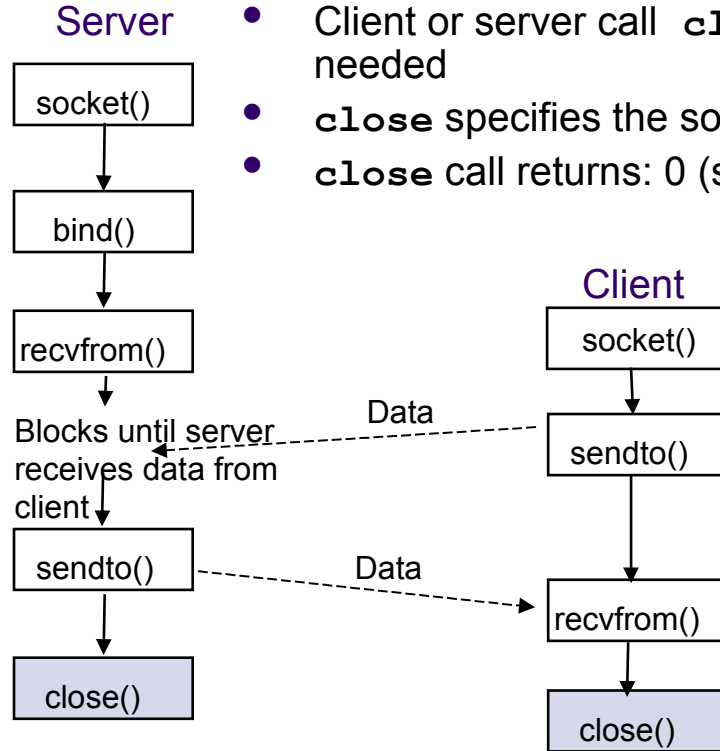


Note: **recvfrom** returns data from at most one **send**, i.e. from one datagram



Socket Close

- Client or server call `close` when socket is no longer needed
- `close` specifies the socket descriptor
- `close` call returns: 0 (success); or -1 (failure)



Socket Calls for Connection-Less Mode



Example-I: TCP Echo Server

- As illustration of the use of system calls and functions, let's see two programs communicate via TCP.
- The client prompts a user to type a line of text and sends it to the server, and reads the data back from the server.
- The server acts as a simple echo server.
- In this example, each program expects a fixed number of bytes from the other end, defined by BUFLen.
- The example code is given in the Textbook Chapter 2.4



TCP Echo Server - Binding

```
/* Bind an address to the socket */
bzero((char *)&server, sizeof(struct
sockaddr_in));
server.sin_family = AF_INET;
server.sin_port = htons(port);
server.sin_addr.s_addr = htonl(INADDR_ANY);
if (bind(sd, (struct sockaddr *)&server,
sizeof(server)) == -1) {
    fprintf(stderr, "Can't bind name to
socket\n");
    exit(1);
}
```

TCP Echo Server - Connections



```
/* queue up to 5 connect requests */  
listen(sd, 5);  
  
while (1) {  
    client_len = sizeof(client);  
    if ((new_sd = accept(sd, (struct sockaddr *)&client,  
&client_len)) == -1) {  
        fprintf(stderr, "Can't accept client\n");  
        exit(1);  
    }  
}
```

TCP Echo Server – Repeated Byte Reads



```
/* Repeated calls to read until all data received */
bp = buf;
bytes_to_read = BUFLLEN;
while ((n = read(new_sd, bp, bytes_to_read)) > 0) {
    bp += n;
    bytes_to_read -= n;
}
printf("Rec'd: %s\n", buf);

write(new_sd, buf, BUFLLEN);
printf("Sent: %s\n", buf);
close(new_sd);
```

TCP Echo Client – Name-to-Address



```
bzero((char *)&server, sizeof(struct sockaddr_in));
server.sin_family = AF_INET;
server.sin_port = htons(port);
if ((hp = gethostbyname(host)) == NULL) {
    fprintf(stderr, "Can't get server's address\n");
    exit(1);
}
bcopy(hp->h_addr, (char *)&server.sin_addr, hp->h_length);
```

TCP Echo Client - Connection



```
/* Connecting to the server */
if (connect(sd, (struct sockaddr *)
&server, sizeof(server)) == -1) {
    fprintf(stderr, "Can't connect\n");
    exit(1);
}
printf("Connected: server's address is
%s\n", hp->h_name);
```

TCP Echo Client – Repeated reads



```
printf("Receive:\n");  
bp = rbuf;  
bytes_to_read = BUFLLEN;  
while ((n = read(sd, bp, bytes_to_read)) > 0) {  
    bp += n;  
    bytes_to_read -= n;  
}  
printf("%s\n", rbuf);
```

Example-II: UDP Echo Server



```
while (1) {
    client_len = sizeof(client);
    if ((n = recvfrom(sd, buf, MAXLEN, 0, (struct
sockaddr *)&client, &client_len)) < 0) {
        fprintf(stderr, "Can't receive datagram\n");
        exit(1);
    }
    if (sendto(sd, buf, n, 0, (struct sockaddr
*)&client, client_len) != n) {
        fprintf(stderr, "Can't send datagram\n");
        exit(1);
    }
}
```


Example: UDP Echo Client



```
gettimeofday(&start, NULL); /*start delay measurement*/
server_len = sizeof(server);
if (sendto(sd, sbuf, data_size, 0, (struct sockaddr *)
    &server, server_len) == -1) {
    fprintf(stderr, "sendto error\n")
    exit(1);
}
if (recvfrom(sd, rbuf, MAXLEN, 0, (struct sockaddr *)
    &server, &server_len) < 0) {
    fprintf(stderr, "recvfrom error\n");
    exit(1);
}
gettimeofday(&end, NULL); /* end delay measurement */
```



Summary: UDP Rliability

- As UDP is unreliable, users may have to take care of reliability assurance by themselves.
- LAN vs. WAN
- Timeout mechanism avoids forever wait
- Re-transmission to get a lost message
- Reordering and de-duplication are required for reliability

